

Problem Statement:

Implement the concept given in below blog link and share the complete steps along with screenshots.

<https://acadgild.com/blog/loading-data-into-hbase-using-pig-scripts/>

Solution:

This document contains the step to execute a **Loading/Transferring of Data into HBase using Pig Scripts**. To achieve this Hadoop cluster is needed with Pig and HBase running on it.

In this example, [Hadoop version: hadoop-2.6.0, Hbase version: hbase-0.98.4-hadoop2-bin and Pig version: pig-0.14.0](#) is used.

The student.txt data set has below column names

Columns Names: StudentName, sector, DOB, qualification, score, state, randomName.

- The 'student.txt' data set will be copied in to HDFS which will be further loaded into HBase.
- Setting the Class Path for Pig

```
[acadgild@localhost ~]$ PIG CLASSPATH=/usr/local/hbase/lib/hbase-server-0.98.14-hadoop2:/usr/local/hbase/lib/hbase-*.jar;
```

- Starting HBase shell and create a table

```
hbase(main):005:0> create 'student','studentdoc'
0 row(s) in 0.8140 seconds

=> Hbase::Table - student
hbase(main):006:0>
```

- Coming out from HBase by typing exit and switch to PIG grunt shell. Loading data from HDFS to Alias relation. Transferring the data inside HBase by STORE command.

```
grunt> studemo = LOAD '/home/acadgild/Downloads/student.txt' USING PigStorage(',') AS(StudentName:chararray,sector:chararray,DOB:chararray,qualification:chararray,score:int,state:chararray,randomName:chararray);
2017-11-17 17:23:22,725 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapreduce.job.counters.limit is deprecated. Instead, use mapreduce.job.counters.max
2017-11-17 17:23:22,727 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum
2017-11-17 17:23:22,728 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt> STORE studemo INTO 'hbase://student' USING org.apache.pig.backend.hadoop.hbase.HBaseStorage('studentdoc:StudentName,studentdoc:sector,studentdoc:DOB,studentdoc:qualification,studentdoc:score,studentdoc:state,studentdoc:randomName');
2017-11-17 17:23:37,360 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapreduce.job.counters.limit is deprecated. Instead, use mapreduce.job.counters.max
2017-11-17 17:23:37,366 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum
2017-11-17 17:23:37,366 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
```

- Once the success message comes , it is confirmed data is loaded inside HBase.

```

Success!

Job Stats (time in seconds):
JobId  Maps    Reduces MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime  MaxReduceTime  MinReduceTime  AvgRe
duceTime  MedianReduceTime  Alias  Feature  Outputs  n/a  n/a  n/a  n/a  0  0  0  0  studemo MAP_0
job_local1068106496_0004  1  0  n/a  n/a  n/a  n/a  0  0  0  0  studemo MAP_0
NLY  hbase://student,

Input(s):
Successfully read 470 records from: "/home/acadgild/Downloads/student.txt"

Output(s):
Successfully stored 470 records in: "hbase://student"

Counters:
Total records written : 470
Total bytes written : 0
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_local1068106496_0004

2017-11-17 17:23:49,029 [main] INFO org.apache.hadoop.metrics.jvm.JvmMetrics - Cannot initialize JVM Metrics with processNam
e=JobTracker, sessionId= - already initialized
2017-11-17 17:23:49,047 [main] INFO org.apache.hadoop.metrics.jvm.JvmMetrics - Cannot initialize JVM Metrics with processNam
e=JobTracker, sessionId= - already initialized
2017-11-17 17:23:49,065 [main] INFO org.apache.hadoop.metrics.jvm.JvmMetrics - Cannot initialize JVM Metrics with processNam
e=JobTracker, sessionId= - already initialized
2017-11-17 17:23:49,198 [main] WARN org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Encount
ered Warning FIELD DISCARDED TYPE CONVERSION FAILED 1 time(s).
2017-11-17 17:23:49,198 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success
!
grunt>

```

- CHECKING THE DATA LOADED IN HBASE

```

hbase(main):002:0> count 'student'
230 row(s) in 0.5060 seconds

=> 230
hbase(main):003:0> scan 'student'

```

```

=====
WILLIAM
WILLIAM
WILSON
WILSON
WILSON
WILSON
WILSON
WILSON
WITMORE
WITMORE
WITMORE
WITMORE
WITMORE
WOODS
WOODS
WOODS
WOODS
WOODS
WORK
WORK
WORK
WORK
WORK
YOUNG
YOUNG
YOUNG
YOUNG
YOUNG
230 row(s) in 6.2130 seconds

```

```

hbase(main):004:0>

```